

Github Link of the project: <https://github.com/sunnysavita10/Finetuning-on-aws>

## Step 1: Create IAM Roles

We'll create two roles:

- SageMakerLLMRole
- LambdaInvokeLLMRole

You can write custom name

### A. Create SageMaker role

- a. Go to AWS Console → IAM → Roles → Create Role
- b. Choose AWS Service → SageMaker
- c. Attach policies:

AmazonS3FullAccess

AmazonSageMakerFullAccess

CloudWatchFullAccess

### B. Create Lambda role

- a. Go to AWS Console → IAM → Roles → Create Role
- b. Choose AWS Service → Lambda
- c. Attach policies:

AmazonSageMakerFullAccess

AmazonDynamoDBFullAccess

AmazonS3FullAccess

CloudWatchLogsFullAccess

## **STEP 2 — Create S3 Buckets for Dataset & Model Artifacts**

1. First bucket for dataset storage
2. Second bucket for model artifacts (after training, SageMaker will save the model)
1. AWS Console → S3 → “Create bucket”
2. Bucket name: llm-finetune-dataset-<yourname>
3. Region = same as your SageMaker region (e.g., ap-south-1)
4. Create a bucket

Why have we created this bucket for keeping the training files

- B. Create Second Bucket (for model artifacts)
1. AWS Console → S3 → “Create bucket”
  2. Bucket name: llm-model-artifacts-<yourname>
  3. Region = same as above
  4. Create a bucket

C. Verify and Note the paths

s3://llm-finetune-dataset-sunny/datasets/

s3://llm-model-artifacts-sunny/models/

## STEP 3 — Create SageMaker Notebook Instance

The screenshot shows the Amazon SageMaker AI Dashboard. At the top left, there's a navigation bar with a menu icon, the text "Amazon SageMaker AI", and a "Dashboard" button with a "New" badge. Below this, there are links for "Getting started", "Dashboard" (which is the current page), and "What's new" (with a "34" badge). On the left side, there are two main sections: "Applications and IDEs" and "Admin configurations". Under "Applications and IDEs", the "Notebooks" link is highlighted with a red box. Under "Admin configurations", there are links for "Domains", "Role manager", "Images", and "Lifecycle configurations". The main content area on the right is titled "Dashboard" and contains several cards: "Amazon SageMaker Studio" (with a brief description), "Environment overview and launch" (with a sub-section for "Filters"), "All active resources" (with a sub-section for "Data prep" and "Ground truth jobs"), and a small "Data prep" card at the bottom.

Getting started

Dashboard New

What's new 34

▼ Applications and IDEs

- Studio
- Canvas
- RStudio
- Notebooks
- Partner AI Apps

▼ Admin configurations

- Domains
- Role manager
- Images
- Lifecycle configurations

Amazon SageMaker Studio

Explore Amazon SageMaker Studio, a machine learning environment where you can build, train, track experiments, deploy models in production, and more.

## Dashboard

This dashboard provides a comprehensive view of your SageMaker resources and environment.

► Environment overview and launch

Overview of the total SageMaker domains, user profiles, and active resources.

► Filters

## All active resources

A current view of your SageMaker AI resources.

## Data prep

Ground truth jobs

Amazon SageMaker AI > Notebook instances > Create notebook instance

### Create notebook instance

Amazon SageMaker provides pre-built fully managed notebook instances that run Jupyter notebooks. The notebook instances include example code for common model training and hosting exercises. [Learn more](#)

**Notebook instance settings**

Notebook instance name: **llm-finetune-notebook**  
Maximum of 63 alphanumeric characters. Can include hyphens (-), but not spaces. Must be unique within your account in an AWS Region.

Notebook instance type: **m1.t5.medium**

Platform identifier: [Learn more](#)  
**Amazon Linux 2, Jupyter Lab 4**

**Additional configuration**

**Permissions and encryption**

IAM role: **SageMakerLLMRole**  
[Create role using the role creation wizard](#)

Enable - Give users root access to the notebook  
 Disable - Don't give users root access to the notebook  
Lifecycle configurations always have root access

Encryption key - optional  
Encrypt your notebook data. Choose an existing KMS key or enter a key's ARN.

No Custom Encryption

**Network - optional**

**Git repositories - optional**

**Tags - optional**

[Cancel](#) [Create notebook instance](#)

Amazon SageMaker AI > Notebook instances > llm-finetune-notebook

### llm-finetune-notebook

[Delete](#) [Stop](#) [Open Jupyter](#) [Open JupyterLab](#) [Edit](#)

**Notebook instance settings**

Name: llm-finetune-notebook	Status: <b>InService</b>	Notebook instance type: m1.t5.medium	Platform identifier: Amazon Linux 2, Jupyter Lab 4 (notebook-al2-v3)
ARN: arn:aws:sagemaker:ap-south-1:730335253621:notebook-instance/llm-finetune-notebook	Creation time: Nov 09, 2025 09:58 UTC	Volume Size: 5GB EBS	Minimum IMDS Version: 2
Lifecycle configuration:		-	

**Git repositories**

Name	Repository URL	Type
There are currently no resources.		

**Permissions and encryption**


IAM role ARN: arn:aws:iam::730335253621:role/SageMakerLLMRole	Root access: Enabled	Encryption key:
---	----------------------	-----------------

**Network**

No custom VPC settings applied

You are screen sharing [Stop share](#)

© 2025, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)



**Keep the necessary requirements in the requirements.txt  
Then install with the below command:  
Pip install -r requirements.txt**

**Note: please check with your requirements.txt path**

1. ls
2. pwd
3. cd /home
4. ls
5. cd ec2-user/
6. ls
7. cd SageMaker/
8. ls
9. cd Finetuning-on-aws/
10. ls
11. git status
12. Clear
13. Activate the env:  
`source /home/ec2-user/anaconda3/bin/activate python3`
14. pip install --upgrade pip setuptools wheel
15. pip install --upgrade cmake
16. sudo yum install -y gcc gcc-c++ make
17. pip install cmake==3.27.0 pyarrow==16.1.0 --no-cache-dir
18. pip install -r ./finetuning/requirements.txt
19. pip cache purge

Open the page below to increase the quota

Quota name	Applied account-level quota value	AWS default quota value	Utilization	Adjustability
ml.g5.xlarge for cluster usage	0	0	0	Account level
ml.g5.xlarge for endpoint usage	0	0	0	Account level
ml.g5.xlarge for notebook instance usage	0	0	0	Account level
ml.g5.xlarge for processing job usage	0	0	0	Account level
ml.g5.xlarge for spot training job usage	0	0	0	Account level
<b>ml.g5.xlarge for training job usage</b>	<b>1</b>	0	0	Account level
ml.g5.xlarge for training warm pool usage	0	0	0	Account level
ml.g5.xlarge for transform job usage	0	0	0	Account level
Studio CodeEditor Apps running on ml.g5.xlarge instances	0	0	0	Account level
Studio JupyterLab Apps running on ml.g5.xlarge instances	0	0	0	Account level
Studio KernelGateway Apps running on ml.g5.xlarge instance	0	0	0	Account level

Create AWS Lambda Function


Go to **Lambda** → **Create function**

**Name:** invoke-llm-lambda

**Runtime:** Python 3.10

**Select Existing Role:** LambdaInvokeLLMRole

## Add the env variable in the lambda



## Paste handler:

```
import json, os, time, boto3
```

```
runtime = boto3.client("sagemaker-runtime")
dynamo = boto3.resource("dynamodb").Table(os.environ["LOG_TABLE"])
ENDPOINT = os.environ["SAGEMAKER_ENDPOINT"]
```

```
def lambda_handler(event, context):
    body = json.loads(event.get("body", "{}"))
    text = body.get("inputs", "")
```

```
    resp = runtime.invoke_endpoint(
        EndpointName=ENDPOINT,
```

```

        ContentType="application/json",
        Body=json.dumps({"inputs": text})
    )
    result = json.loads(resp["Body"].read().decode())

    log_item = {
        "id": f'{int(time.time())#{context.aws_request_id}}',
        "prompt": text,
        "response": result,
        "timestamp": int(time.time())
    }
    dynamo.put_item(Item=log_item)

    return {
        "statusCode": 200,
        "body": json.dumps({"result": result})
    }
}

```

## Create DynamoDB Table (for logs)

Go to DynamoDB → Create table

- Table name: llm\_queries
- Partition key: id (String)
- Sort key: (optional)

## Deploy-test with the sample

```
{
    "body": "{\"inputs\": \"Summarize the text about heart disease\"}"
}
```

## Create API Gateway (REST)

1. Go to API Gateway → Create API → REST API → Build
  - Name: LLMInferenceAPI
2. Create resource /predict
3. Create method POST
  - Integration type: Lambda Function
  - Choose invoke-l1m-lambda
4. Deploy API → stage: prod

**Now create a UI for the inferencing and RAG appp**

**Create a virtual env and install the  
requirements\_inference.txt**

1. uv python list
2. uv venv env --python cpython-3.11.13-windows-x86\_64-none
3. env/Scripts/activate
4. source env/Scripts/activate→ for bash terminal
5. uv pip install -r requirements\_inference.txt

Then after creating a app will run it via these below commands

6. streamlit run inference\_app.py